

Input-Output Management

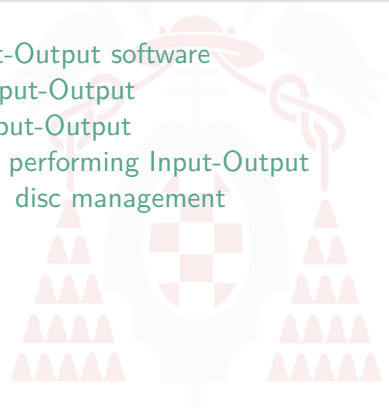
E. Campo M. Knoblauch Ó. López J. Clemente

Departamento de Automática
Universidad de Alcalá



/gso>


Index

- 1 Introduction
 - 2 Layers of Input-Output software
 - 3 Elements of Input-Output
 - 4 Mapping of Input-Output
 - 5 Techniques for performing Input-Output
 - 6 Sample device: disc management
- 

Issues with I/O

- Extensive use of I/O
 - More I/O concurrence → better usage of the system
- Every I/O device has its own idiosyncrasy
 - Diverse functionality
 - Storage (Discs)
 - User interface (Keyboard, mouse)
 - Communications (Network cards, modem)
 - Etc.
 - Ways to access them
 - Different degrees of autonomy and intelligence
- Devices need help from the kernel
- Different speed

Speeds of I/O devices



I/O management

- The operating system is the link between the user application and the I/O hardware
- Goals of the I/O subsystem of the O.S.:
 - Hide HW from applications
 - Uniform access interface (uniform names)
 - Device-independence
 - Ability to write programs that access I/O devices without knowing in advance the type of device
 - Error handling
 - Management of different device types
 - Shared or dedicated
 - Character mode or block mode
 - Sequential access or random access
 - Etc. . .

Layers of I/O software (1/2)

I/O calls, I/O
formatting, spooling

User space I/O

Naming, protection,
locking, buffering,
device assignment

Device independent I/O

Value assignment to
device registers,
status checks

Disc
handler

Serial line
handler

Sound
handler

Graphics
handler

Keyboard
handler

Mouse
handler

INT management

Interrupt handler

Execution of I/O
operation

Disc
controller

Serial line
controller

Sound
controller

Graphics
controller

Keyboard
controller

Mouse
controller




Layers of I/O software (2/2)

- I/O hardware can be updated and modified:
 - The O.S. provides mechanisms for updates
 - What happens if a new HW element is introduced?
 - Monolithic approach versus microkernel approach
- The I/O interface with the final user:
 - Uniform access
 - All devices look the same
 - Device-independent
 - Hides HW from upper levels
 - The user doesn't bother accessing every device
 - Same access primitives, regardless of device type

Elements of I/O

- Every I/O element is directly or indirectly connected to the system bus
- Usually divided in two parts:
 - Device controller or adapter
 - Printed circuit card, connected to an expansion slot of the computer
 - Has some state, control and data registers
 - The device itself
 - Connected to the controller through a slot, wire or similar, which can be standard (i.e. IDE or SCSI)



Interrupt Handlers

- Interrupts are inevitable in most I/O operations
 - They should be hidden as much as possible
 - The O.S. can/must carry out several tasks:
 - Store registers and PSW, and prepare a context and a stack in order to execute an ISR
 - Send an acknowledgement to the interrupt controller
 - Execute the ISR, which will interact with the controller that generated the interrupt
 - Sometimes, wake up a process implied in the I/O and invoke the scheduler
 - Prepare the context and recover the registers of the new process
 - Return from the interrupt and execute the new process

Device Handlers

- They contain specific code for every device
- Duties:
 - Initialize the device
 - Accept abstract read/write requests from the device-independent software
 - Record events
- Usually written by the manufacturer of each hardware device
 - Important: O.S. designers must provide a well defined model of what a device handler does and how should it interact with the O.S.
 - Standard interface for block mode devices and character mode devices

Mapping of Input-Output

- Communication between controller and CPU
 - Control registers: The CPU can read/write them to command the device, find out its state, etc.
 - Data buffers where the O.S. can read/write
- How?
 - Assign an I/O port to every register
 - Port: 8 or 16 bits integer
 - Specific instructions for reading/writing (IN/OUT)
 - Establish a correspondence between registers or memory positions
 - Accesses to these positions actually access the registers
 - Standard load/store instructions (i.e. MOV)
 - Hybrid
 - Ports for the registers, memory for the buffers


Techniques for performing I/O (1/2)

- Peripheral's operations are asynchronous with respect to those of the processor
- Different communication mechanisms:
 - Polling
 - Interrupts
 - Direct memory access

Techniques for performing I/O (2/2)

- Programmed I/O (Polling)
 - The CPU waits until the device is ready to send the next character (polling)
 - Simple, but with busy wait
- I/O by means of interrupts
 - The CPU starts the I/O operation, and it continues with an interrupt
 - The process that makes the I/O call gets blocked, thus avoiding busy wait
 - An interrupt is produced for every character
- I/O by means of DMA
 - It is the DMA controller who serves characters
 - The DMA controller basically carries out programmed I/O
 - The process that makes the I/O call gets blocked
 - Only one interrupt is produced, and busy wait is avoided

Hardware of a disc



- Disc divided in 512-bytes sectors, referenced by:
 - Cylinder, Head, Sector (CHS)
- Access times
 - Seek time
 - Rotational latency
 - Transference time
- Which one can be optimized?
- On modern discs, the internal structure is transparent. Sectors are referenced with Logic Block Addressing (LBA)

Disc handler

- Goal: read and write disc sectors
- Basic procedure:
 - Initialize DMA
 - Start the motor (floppy discs or hard discs with power saving)
 - Move the heads to the right position
 - Read or write data (completion is signalled by the controller with an interrupt)
 - Stop the motor (floppy discs)
- The operations to be carried out will depend on the degree of intelligence and autonomy of the disc controller

Scheduling of disc I/O (1/2)

- There are several scheduling possibilities
- In order to evaluate every algorithm, a list of requests is used:
 - 98, 183, 37, 122, 14, 124, 65 y 67 (with head starting at track 53)
- First Come - First Served (FCFS)
 - Simplest method to schedule and implement
 - Acceptable with small workloads (equitable)
 - Traverses 640 tracks

Scheduling of disc I/O (2/2)

- Shortest Seek Time First (SSTF)
 - Minimizes seek times
 - Response times are not equitable, external tracks are discriminated and might suffer starvation
 - Traverses 236 tracks
- SCAN
 - Also known as the “elevator algorithm”
 - The C-SCAN variant provides more uniform response times
 - Traverses 236 tracks
- Nowadays, it is the HW disc controller who optimizes the disc I/O requests

Bibliographic references

-  [Sánchez, 2005] S. Sánchez Prieto.
Sistemas Operativos.
Servicio de Publicaciones de la UA, 2005.
-  [Tanenbaum, 2009] A. Tanenbaum.
Sistemas Operativos Modernos.
Ed. Pearson Education, 2009.
-  [Stallings, 1999] W. Stallings.
Organización y arquitectura de Computadores.
Ed. Prentice Hall, 1999.
-  [Silberschatz, 2006] A. Silberschatz, P. B. Galván y G. Gagne
Fundamentos de Sistemas Operativos.
McGraw Hill. 2006