

# Memory Management

E. Campo   M. Knoblauch   Ó. López   J. Clemente

*Departamento de Automática*

Universidad de Alcalá



/gso>

# Index

## 1 Introduction: from programs to processes

- Format of an executable file
- Memory map of a process
- Example of program and process
- Addressing abstractions

## 2 Principles of memory management

- Memory hierarchy in a computer
- Principle of locality
- Fragmentation
- Relocation
- Protection and sharing

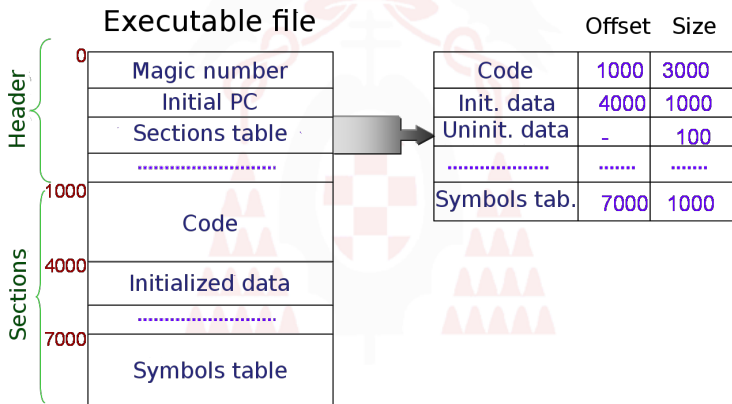
## 3 Memory management mechanisms

- Historic evolution
- Non-contiguous partitioned memory
- Segmentation
- Paging
- Paged segmentation

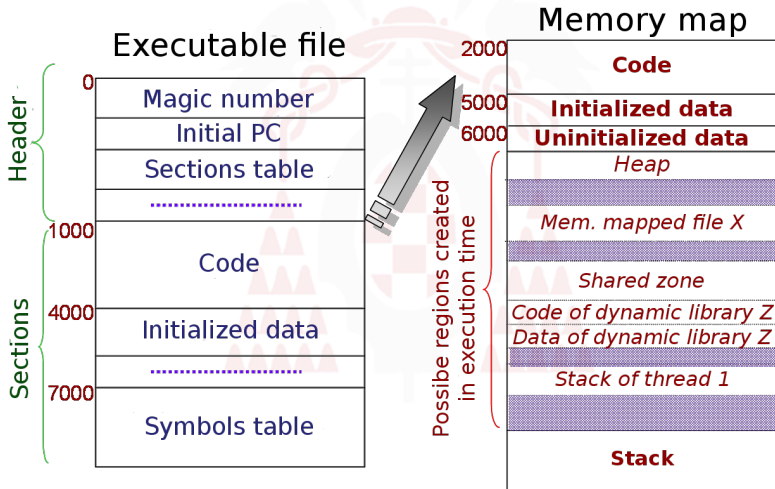
## 4 Case studies

- Pentium's memory management

# Format of an executable file



# Memory map of a process



# Program and process

```
...  
char * progame, int cont = 1;  
  
void Func (int x)  
{  
    int result = 0;  
    char character = 'a';  
  
    if (cont)  
        cont = result ++;  
    ...  
    return;  
}  
  
main (int argc, char * argv[])  
{  
    int i;  
    char * progame;  
  
    Func (cont);  
    progame = (char *) malloc (1 + strlen(argv[0]));  
    ...  
    free (progame);  
    ...  
    exit (0);  
}
```


# Addressing abstractions

## Address space

Set of referentiable addresses

- Virtual address space  $\Rightarrow$  independent for every process
- Physical address space  $\Rightarrow$  shared amongst all processes
- Processes only reference virtual addresses
- There must be a translation from virtual to physical address, transparent to the process

# Addressing abstractions




# Memory hierarchy in a computer

- The organization of memory into a hierarchy is an attempt to enhance computers' performance
- Based on: programs' locality + technological advances in memory designs
- Fast memories: small capacity, expensive
- Slow memories: large capacity, cheap



# Scheme of memory hierarchy



# Principle of locality

- During time intervals, processes tend to concentrate references in a subset of their address space

Donald Knuth [1971]:

Programs typically have a very jagged profile, with a few sharp peaks. [..]  
We also found that less than 4 per cent of a program generally accounts for more than half of its running time.

- It's an empirical property
- There are two types of locality:
  - Spatial locality
  - Temporal locality

# Spatial locality

- Once a memory position has been referenced, odds are that near positions will be referenced either.
- Supporting this remark:
  - Sequential execution of code
  - Programmers' tendency to put related variables together
  - Access to data structures like stacks or arrays

# Temporal locality

- Once a memory position has been referenced at instant  $t$ , odds are that it will be referenced again at instant  $t + \Delta t$
- Supporting this remark:
  - Loops
  - Subroutines
  - Stacks

# Fragmentation

## Fragmentation

Waste (inefficient use) of the available free memory due to the management mechanism employed

- There are two types: internal and external
- Internal fragmentation
  - Caused by the difference of size between the memory partition and the object allocated inside it
- External fragmentation
  - Caused by the inability to use memory between partitions

# Relocation

## Relocation

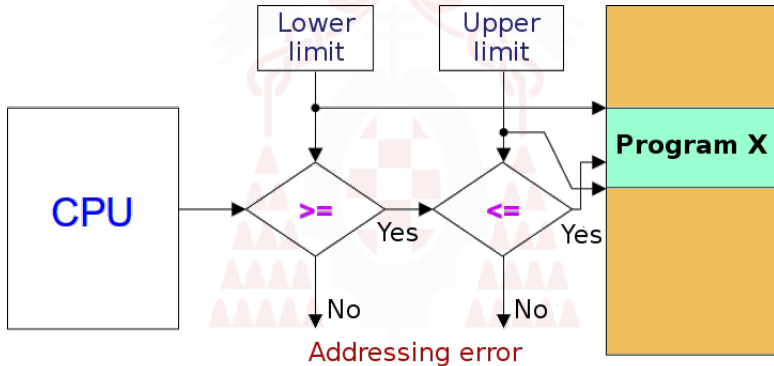
Assignment of addresses to the different parts of a program (code, data, stack...)

- Depending on **when** is fixed the final location, relocation will happen in the compilation stage, in the loading stage, or during the execution stage
- Static relocation
  - Carried out before or while loading the program
  - Once started, programs cannot be moved
- Dynamic relocation
  - The translation from virtual address to real (physical) address is carried out in execution time
  - It requires additional hardware (MMU)
  - Programs can be moved in execution time

# Protection and sharing

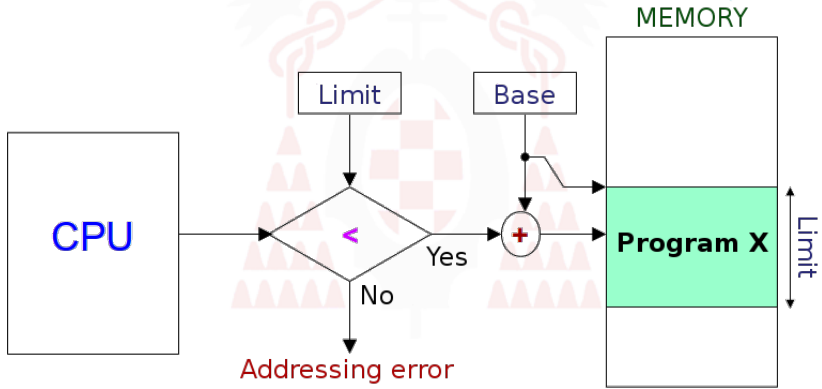
- Need to restrict memory access:
  - Operating system
  - User processes
- Protection methods
  - Limit registers
  - Base and limit registers
  - Memory protection bits
  - Access rights in translation tables
  - Where are they stored?
- How to share memory between processes?

# Limit registers





## Base and limit registers



# Historic evolution

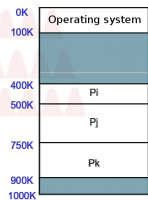
- Bare machine
  - The system provides no service
- Monolithic monitor
  - In addition to the operating system, there is just one process
- Contiguous partitioned memory
  - Multiprogramming with a fixed number of tasks (MFT)
    - Fixed-size partitions
    - Created while booting the system
  - Multiprogramming with a variable number of tasks (MVT)
    - Variable-size partitions
    - Created when processes need them
- Non-contiguous partitioned memory

# Non-contiguous partitioned memory

- The contents of a process can be distributed in separated memory partitions
- The memory is organized in partitions:
  - Variable size  $\Rightarrow$  segments
  - Fixed size  $\Rightarrow$  frames

## Partitions description table

- Independent for every process
- Built when the process is loaded



| Partition number | Base | Size | State    |
|------------------|------|------|----------|
| 0                | 0K   | 100K | ASSIGNED |
| 1                | 100K | 300K | FREE     |
| 2                | 400K | 100K | ASSIGNED |
| 3                | 500K | 250K | ASSIGNED |
| 4                | 750K | 150K | ASSIGNED |
| 5                | 900K | 100K | FREE     |

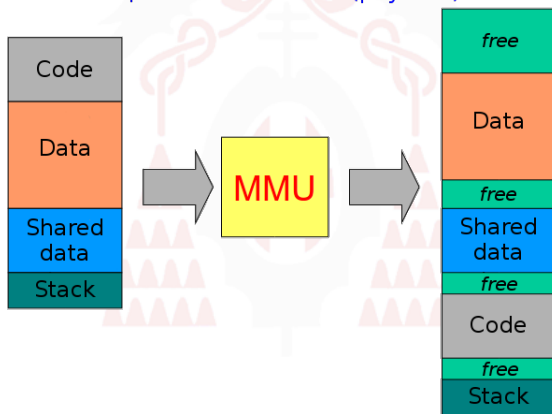
# Segmentation

- The physical memory is initially organised as a unique empty block, where variable-size partitions (segments) are created as required
- The virtual address space is organized in segments
- Includes protection mechanism and allows sharing
- Virtual addresses are composed of two elements: segment number and offset
- The partition table is called Segment Table (ST)
- If the ST is too big, it has to be stored in main memory, pointed by a register (STBPR)  $\Rightarrow$  every access requires two references to memory

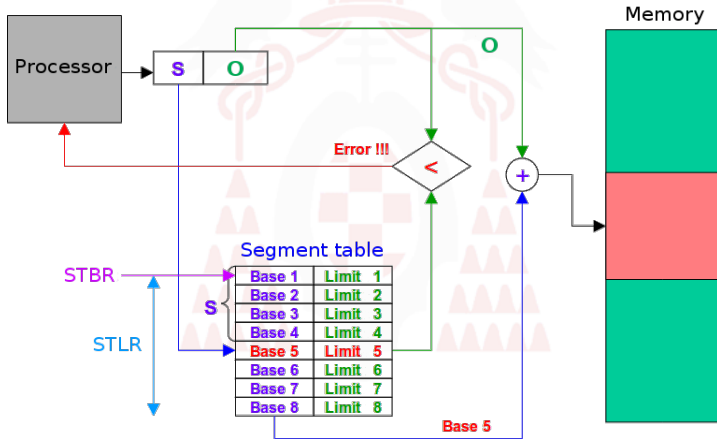
## Segmentation: logical scheme

Virtual address space

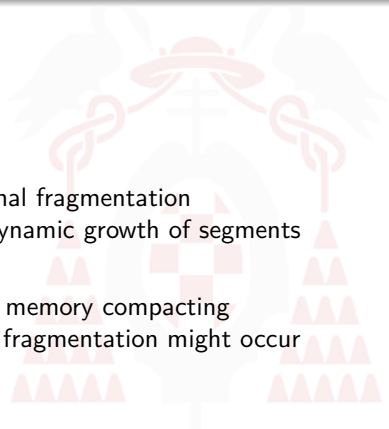
Real (physical) address space



## Segmentation: physical scheme



# Segmentation considerations

- 
- Advantages:
    - No internal fragmentation
    - Allows dynamic growth of segments
  - Drawbacks:
    - Requires memory compacting
    - External fragmentation might occur

# Paging

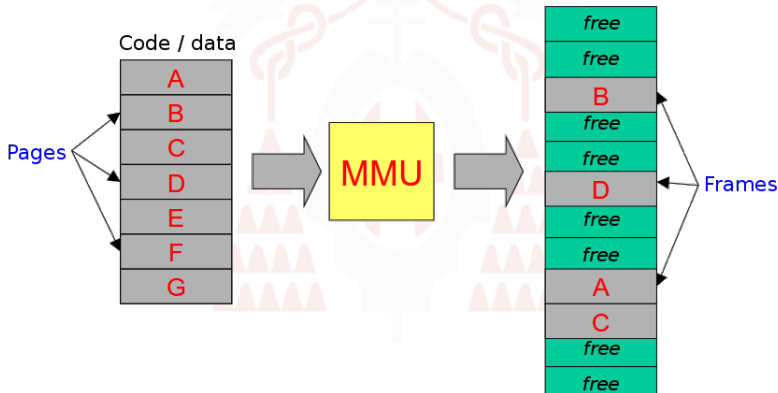
- The physical memory is initially organised in fixed-size partitions (frames)
- The virtual address space of every process is divided in fixed-size blocks (pages)
- Virtual addresses are composed of two elements: virtual page number and offset
- Includes protection mechanism and allows sharing
- The partition table is called Page Map Table (PMT)
- If the PMT is too big, it has to be stored in main memory, pointed by a register (PTBPR)



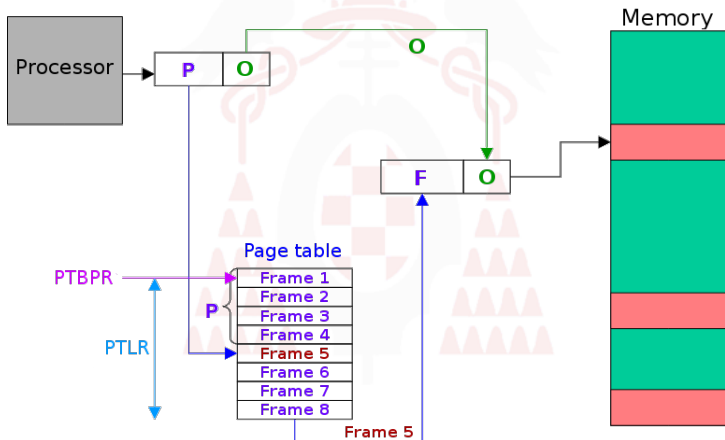
# Paging: logical scheme

Virtual address space

Real (physical) address space



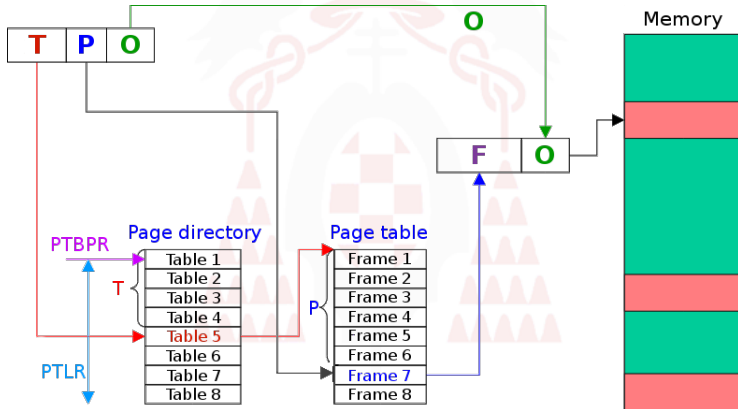
## Paging: physical scheme



# Paging considerations

- Advantages:
  - No external fragmentation
- Drawbacks:
  - Internal fragmentation might occur
- With large pages internal fragmentation grows but the PMT gets smaller, and vice-versa
- If the number of pages is high, the amount of memory occupied by the PMT can be prohibitive. In such cases, the PMT itself has to be paged.

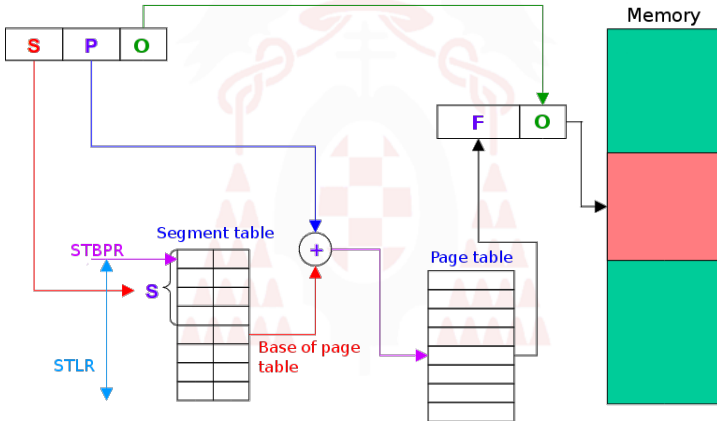
# Paged paging



## Combination of mechanisms


- It is possible to combine the schemes of paging and segmentation
- The advantages of both are obtained, at the price of complicating the hardware
- Possible combinations:
  - Paged segmentation
  - Segmented paging (not used in practice)

# Paged segmentation: logical scheme



# Pentium's MMU

- The Pentium supports segmentation, paging and paged segmentation (the most usual)
- The logical address is composed of a segment selector (13+1 bits) and an offset (32 bits)
- The segment selector is the value contained in one of the next registers: CS, DS, ES, SS, FS, GS




## Format of the segment descriptor

- LDT (Local Descriptor Table)  $\Rightarrow$  one per process
- GDT (Global Descriptor Table)  $\Rightarrow$  one per system
- Maximum number of entries per table  $\Rightarrow 2^{13}$
- Each entry in the segment table is called *descriptor*
- Descriptor size  $\Rightarrow$  8 bytes

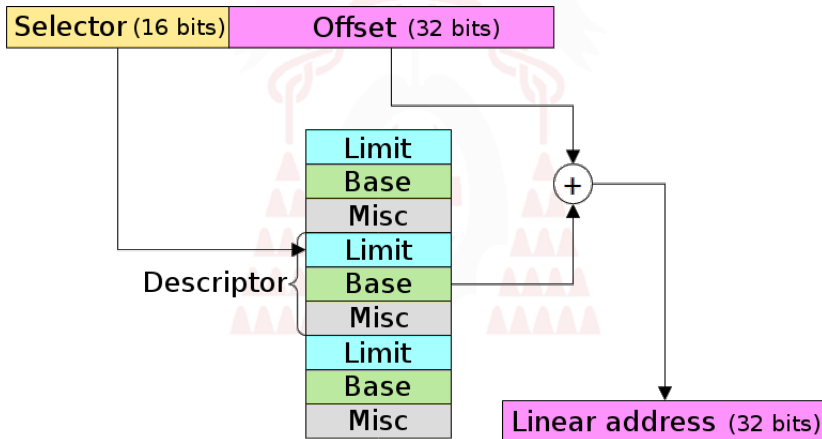
### Segment descriptor

- Base address (32 bits)
- Limit (20 bits)
- Attributes and privileges (12 bits)







## Pentium's segmentation stage



## Pentium's paging stage



# Bibliographic references I

-  [Sánchez, 2005] S. Sánchez Prieto.  
*Sistemas Operativos.*  
Servicio de Publicaciones de la UA, 2005.
-  [Tanenbaum, 2009] A. Tanenbaum.  
*Sistemas Operativos Modernos.*  
Ed. Pearson Education, 2009.
-  [Stallings, 1999] W. Stallings.  
*Organización y arquitectura de Computadores.*  
Ed. Prentice Hall, 1999.
-  [Silberschatz, 2006] A. Silberschatz, P. B. Galván y G. Gagne  
*Fundamentos de Sistemas Operativos.*  
McGraw Hill. 2006

## Bibliographic references II



D.E. Knuth.

An Empirical Study of FORTRAN Programs

*Software—Practice and Experience*, vol. 1, 105–133, 1971.